MS-ConvLSTM: about half the training energy of a same-depth ConvLSTM, with lower error

Bottom-left is best. MAE: mean of the 4 trained models; energy: mean of the runs on sirocco21/24; bars: standard deviation. ConvLSTM 2L is cheaper but clearly less accurate.

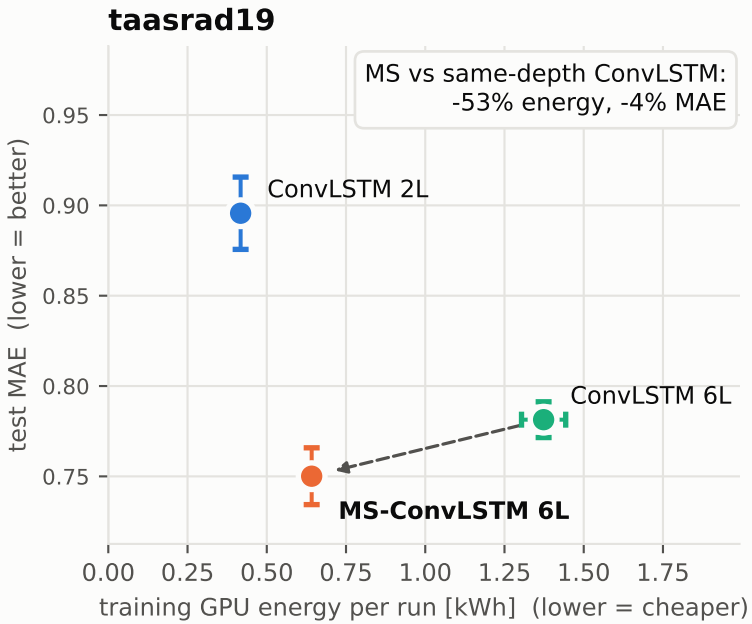

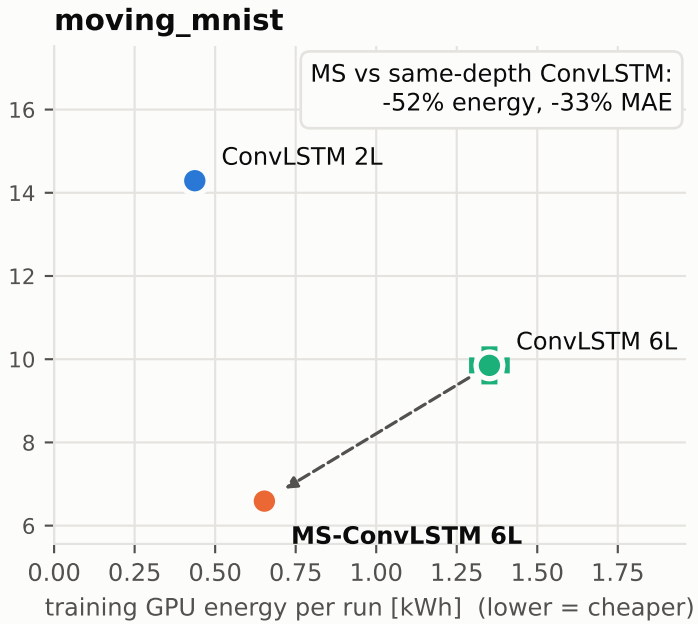

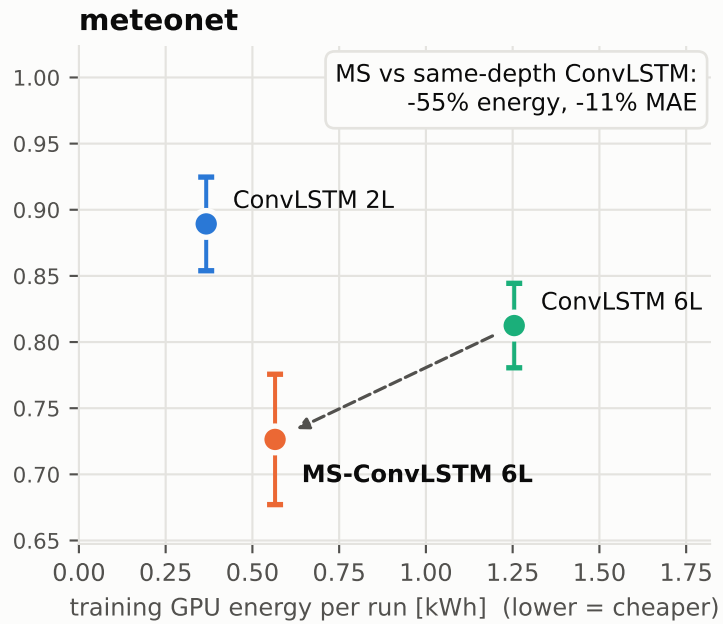

100 epochs, batch 10, 64 channels, NVIDIA A100-PCIE-40GB (plafrim). Energy = active GPU, from IOPS 1 s nvidia-smi sampling, runs on sirocco21/24 only (n = 1 to 4 per point). MAE on the test split, all 4 repetitions.